

Delta Logics: Logics for Change

ANONYMOUS AUTHOR(S)

ACM Reference Format:

Anonymous Author(s). 2018. Delta Logics: Logics for Change. 1, 1 (April 2018), 16 pages. <https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

1 AN EXPLICIT FORMULATION

We describe an explicit formulation that offers more insight into the working of delta-logics. Recall that the main idea is to separate the heap into two parts: one that changes and one that does not and, in correspondence, express the VC as a boolean combination of two kinds of formulae describing each one of those parts. This is a formulation in a static logic with one (copy of) parameterised recursive function. We then use frame reasoning on the unchanging context to relate the two forms of the parameterised recursive function, SMT reasoning on the bounded Δ to relate the two sets of parameters and solve the mutual constraints to check VC validity. In the general formulation the presence of the two different sets of parameters does not necessarily make clear this intuition. However in some cases, and in particular the case of lists and measures described in Section 6, the recursive function R satisfies the following condition for all sets of parameters P :

$$\forall x \exists y, y \in \Delta. R(x, P) = R(x, \perp_R) \oplus_R P(R, y)$$

where $P(R, x)$ is a slight abuse of notation referring to the parameter variable in P ‘corresponding’ to x for R , \perp_R is a constant of the type of $P(R, x)$, and \oplus_R is some binary operator of appropriate signature such that the above equation is valid. In fact, we can make it a little more general and make \perp_R a function of x as well. The \oplus_R operator may also require more than just the corresponding parameter, and may make use of the parameters corresponding to $Rank_R$ as well.

In this case, observe that we can explicitly apply frame reasoning to conclude that the valuation of the term $R(x, \perp_R)$ does not change by pointing out that the underlying heaplet of its definition is exclusively outside Δ , which does not change. This is true for any x (from our formulation of the definition R above). The changed parameter $P'(R, x)$ can then be recombined through the \oplus_R operator with this term (since the equation is valid for all sets of parameters) to yield a valuation for $R(x, P')$.

Example: Consider the example given in Section 2:

$$\{ls(x, nil) \wedge ls(w, nil) \wedge y \notin hls(w, nil)\} y.next := w \{ls(x, nil)\}$$

Even in the simple subcase where $w \notin hls(x, nil)$, we have to argue the validity of the postcondition in both cases, namely $y \in hls(x, nil)$ and $y \notin hls(x, nil)$. As we saw earlier, vanilla frame reasoning will not work for the former case. In the delta-logic formulation, let the set of parameters for the pre- and post- states be P, P' respectively (the individual variables are thought to be primed versions as well). The relativised definition of ls with nil as the terminal point given in Section 6 satisfies the following equation for any h :

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.

© 2018 Association for Computing Machinery.

XXXX-XXXX/2018/4-ART \$15.00

<https://doi.org/10.1145/nnnnnnnn.nnnnnnnn>

$$ls_{nil}^P(h) \iff ls_{nil}^{true}(h) \wedge LS_{nil}^{f(h)}$$

where $\perp_{ls} = true$, $\oplus_{ls} \equiv \wedge$ and we have applied the same abuse of notation in our description to the LS parameters and $f(x)$ is a Skolem function. Observe that in this case, where y and w are Δ , it is possible to easily conclude that LS_{nil}^y holds, and since for none of the locations $d \in \Delta$ the parameter change LS_{nil}^d to $LS_{nil}'^d$ is from true to false (in fact for y it may have been false to true), and that $ls_{nil}^P(x)$ held at the beginning, we have $ls_{nil}^{P'}$ holds on the post-state, which is what we wanted.

Observe here that the application of simple frame reasoning to our formulation yields a more fine-grained frame rule in that we can still conclude the truth of recursive predicates if the parameters they depend on become only ‘more’ true.

This condition is true of some of our list measures as well: with the \perp_R and \oplus_R being respectively: empty set and set-union for *hls*, similarly empty multiset and multiset-union for *mskeys*, and 0 and standard addition for *len*. This formulation illustrates the motivation and working of delta-logics: it is about enabling simpler reasoning by the use of close and fine-grained frame reasoning while being able to handle complex conditions on bounded worlds.

2 INTRODUCTION

Classical logics, such as first-order logic with least fixpoints or higher order logics, are often static in the sense that formulae describe the state of a single world. Verification conditions of imperative programs describe typically at least two different worlds: the pre-state and the post-state of a program. Consequently, expressing validity of verification conditions naturally involves the challenge of expressing the evolving worlds of program states in a static logic. The classical notion of strongest post-condition for programs with scalar variables does precisely this— it expresses the precondition, the intermediate states of the program, and the post-state using *auxilliary* first-order variables that capture the scalar variables at these states. The weakest precondition, again for programs with scalar variables, solves the same problem.

The focus of this paper is in generating logical formulations of verification conditions for program snippets that manipulate the heap. In this setting, the heap consists, minimally, of a set of pointer fields that are modeled by *first-order functions*, and the program’s execution alters these functions. Consequently, the translation of validity of verification conditions to validity of formulae changes considerably.

Let us consider a set of pointer fields \bar{p} and a recursive definition of a unary predicate or function $R(x)$ defined using least fixpoints over \bar{p} . Typical functions include properties such as “ x points to a linked list segment ending at the location z ”, “the length of the list pointed to be x ”, “ x points to a binary search tree”, “the set of keys stored in the tree pointed to by x ”, “the heaplet defined by the tree pointed to by x ”, etc. Consider a Hoare triple of the form $@pre(\bar{x}, \bar{p}, R) S @post(\bar{x}, \bar{p}, R)$: the pre- and post-conditions use the recursively defined predicate/function R .

One simple approach is to capture the precondition using logic and use the *frame rule* to reason soundly (but incompletely) about the post-state— i.e., we can simply ignore the definition of R on the transformed heap and infer that $R'(x)$ holds if it held in the pre-state and the modified portion of the heap did not intersect with the underlying heaplet of $R(x)$ (from the definition of R). This is, in practice, a very *convenient and simple* reasoning that often works, but is incomplete. For example, consider the following Hoare triple with pre/post conditions written in FO+*lfp*, where $list(x)$ means x points to a list, and in which case, $hlist(x)$ stands for the set of locations in that list:

$$\{list(x) \wedge list(w) \wedge y \notin hlist(w)\} y.next := w \{list(x)\}$$

In this case, the Hoare triple does hold, but frame reasoning is insufficient to argue it— since y can be in $hlist(x)$ in the pre-state, the frame rule is insufficient in inferring that the postcondition

continues to hold. (The above can also be formulated in separation logic, of course, and using frame reasoning **alone** will continue to be ineffective.)

Let us now focus on basic blocks that do not involve function calls. We would like to generate precise verification conditions in such cases, while falling back on frame reasoning for function calls. There are several approaches in the literature that argue for this: for example the Grasshopper suite of tools handle such blocks accurately [14], and there is work on expressing weakest pre-conditions in separation logic for such blocks using the magic wand.

One precise formulation of the verification condition is $@pre(\bar{x}, \bar{p}, R) \wedge T(\bar{x}, \bar{x}', \bar{p}, \bar{p}') \wedge @post(\bar{x}', \bar{p}', R')$ where T describes the effect of the program on the stack and the heap, describing how the scalar variables \bar{x} and pointer-fields \bar{p} have evolved to the scalar variables \bar{x}' and \bar{p}' . Most importantly, the new recursive definitions R' are *reformulated* by replacing \bar{p} in the definition of R with \bar{p}' .

Though the above is a precise formulation of the verification condition, it has several drawbacks. First, there is a heaplet H modified by the program, and the formula will have conjuncts of the form $\forall y \notin H. p'(y) = p(y)$, for every $p \in \bar{p}$. This introduces universal quantification, which is harder to reason with automatically. However, for basic blocks that do not involve function calls, i.e., when H is finite, we can map this into a decidable quantifier free logic (modeling p as an array and p' as an update to the array). Second, the new definition of R' depends on \bar{p}' , which in turn depends on the various constraints introduced by the basic block. For example, pointer fields may change depending on complex properties involving the data elements stored in the heap. Reasoning with R' (which involves least fixpoints), coupled with such constraints, is daunting.

Surely, there must be a simpler formulation of the verification condition! Small changes to the heap do cause global changes and can dramatically affect the valuation of recursively defined predicates/functions, which are global. But surely, the effect on the semantics of R (changing into R') must be expressible in a simpler localized fashion.

In this paper, we describe a class of logics, called *delta-logics*, that are precisely meant to address the above. Formulae in delta-logics are Boolean combinations of two distinct kinds of formulae: one kind, called *delta-specific* formulae, strictly talk about the modified portion of the global heap (identified by a bounded set of locations ' Δ ') without using any recursive definitions; while the other kind, called *contextual/context-logic* formulae, strictly talk about the unbounded portion different from Δ using recursive definitions. A set of interface variables are used to communicate information between the Δ and the rest of the heap: its 'context'. In particular, a recursive definition R over the unbounded context is *parameterized* over a set of first-order communication variables P^R , where P^R summarizes the values of R within Δ . These variables themselves can, of course, depend on the value of R outside Δ as well, setting up mutual constraints.

We prove several key results. First, we show a *separability theorem* in Section 4 that shows that any quantifier-free FO formula with recursive definitions (with *lfp* semantics) can be expressed in delta-logics, i.e., as a Boolean combination of delta-specific formulae and contextual formulae. We prove this by setting up communication between the two portions of the heap as described above, and it turns out that a new set of recursive definitions and communication variables involving the *rank* for recursive definitions is needed to accurately capture least fixpoints. These ranks can however be constrained to be *bounded integers* as opposed to ordinals (though the heap is infinite).

Second, we argue that verification conditions of basic blocks without function calls can be expressed precisely using delta-logics. The key idea here is to set up *two parameterized sets* of recursive definitions, expressing the pre-condition using one and the post-condition using the other. The change to the heaplet can be described in quantifier-free FOL as usual, and the pre- and post-conditions can be translated to delta-logic formulae using the separability theorem. Notice that this captures naturally the changing Δ portion of the heap while keeping the context unchanged

and hence removes the need for the universally quantified formula asserting that pointers in the context have not changed.

Then, we turn to specific delta-logics and explore decidability results for them. We define a delta-logic that expresses properties of list segments along with a variety of measures on them, including their heaplets (for expressing separation properties), their lengths, the multisets of keys stored in them, the min/max keys stored in them, and their sortedness. We consider the contextual logic corresponding to this delta-logic, and by exploiting the simplicity of delta-logics (namely, that the recursive definitions change only w.r.t the communication parameters which correspond to base cases of these definitions), we show that they can be transformed to equivalent quantifier-free formulae *without* recursive definitions. This leads us to a decision procedure for delta-logics for linked lists with all the six measures above, and hence a decidable logic for verification conditions of programs manipulating linked lists with pre- and post-conditions expressed using the above measures.

Finally, we implement and evaluate our technique by expressing VCs using delta-logics and validating them using our decision procedure on a suite of programs, and show it to be effective. To the best of our knowledge, this is the most expressive decidable logic over lists in existing literature.

While our decidability result is interesting in its own right, we emphasize the main contribution of this paper is, in our view, the definition and advocacy of delta-logics as a logic for verification conditions. Traditional ways of reducing validation of verification conditions to logic embed the verification in logics that have much more expressive power than necessary, and harder to reason with; we think it is a reduction from an easier problem to a harder problem! Delta-logics explicitly delineate the changed portion of the heap by only allowing recursive definitions to be parameterized over first-order communication variables and avoid introducing quantification, making them easier to reason with.

3 DELTA-LOGICS

In this section, we define delta-logics extending many-sorted first-order logic with least fixpoints and background axiomatizations of some of the sorts.

Let us fix a many-sorted first-order signature $\Sigma = (S, \mathcal{F}, \mathcal{P}, \mathcal{C}, \mathcal{G}, \mathcal{R})$ where $S = \{\sigma_0, \dots, \sigma_n\}$ is a nonempty finite set of sorts, \mathcal{F} , \mathcal{P} , and \mathcal{C} are sets of function symbols, relation symbols, and constant symbols, respectively, and \mathcal{G} and \mathcal{R} are function and relation symbols that will be recursively defined. These symbols have implicitly defined an appropriate arity and a type signature.

Let σ_0 be a special sort that we refer to as the *location* sort, which will model locations on the heap. The other sorts, which we refer to as background sorts, can be arbitrary and constrained to conform to some theory (such as a theory of arithmetic or a theory of sets).

We assume the following restrictions:

- We assume all functions in \mathcal{F} map either from tuples of one sort to itself or from the foreground sort σ_0 to a background sort σ_i . Relations in \mathcal{P} are over tuples of one sort only.
- The functions in \mathcal{F} whose domain is over the foreground sort σ_0 are *unary*. Also, relations over the foreground sort σ_0 are unary relations.
- Recursively defined functions (in \mathcal{G}) are all unary functions from the foreground sort σ_0 to the foreground sort or a background sort. Recursively defined relations (in \mathcal{R}) are all unary relations on the foreground sort σ_0 .

The restriction to have unary functions from the foreground sort (which models locations) is sufficient to model pointers on the heap (unary functions from σ_0 to σ_0) and to model data stored in the heap (like the key stored at locations modeled as a function from σ_0 to a background sort of integers). This restriction will greatly simplify the presentation of delta-logics below. The restriction

of having unary recursively defined functions and relations will also simplify the notation. Note that recursive definitions such as $lseg(x, y)$ that are binary can be written recursively as unary relations such as $lseg_y(x)$ (i.e., parameterized over the variable y) with recursion on the variable x .

The logic $FO+lf\!p$ that we use consists of a set of recursive definitions of (unary) predicates and functions (with least fixpoint semantics) and a quantifier-free formula that uses these definitions. For a simpler exposition, a definition of a recursive function R will be of the form: $R ::=_{lf\!p} \varphi(x, R(p_1(x)), \dots, R(p_n(x)))$ where p_i are unary functions and φ is monotonic, i.e., the $R(p_i(x))$ occur in the definition under an even number of negations.

The lattice that we have for computing the least fixpoints for predicates is $\{\top, \perp\}$ with $\perp < \top$. For functions whose range is a domain \mathcal{D} , this lattice is $\mathcal{D} \cup \{\perp\}$, where $\perp < d$ for every d in \mathcal{D} . The semantics of atomic formulas is that they evaluate to *false* whenever they involve a term involving \perp , when the formula is under an even number of negations; and to *true* otherwise.

Furthermore, we restrict the kind of recursively defined predicates to have definitions of the form $R(x) ::= \varphi(x, e)$ where e is a set of conjunctions involving $R(p_i(x))$. This restricts definitions to have a unique dependence on their successors, and most natural definitions for heaps satisfy this property. Similarly, the restriction on recursively defined functions is that they have definitions of the form $G(x) ::= t(x, e)$ for some term t , such that e is a term involving terms of the form $G(p_i(x))$.

We parameterize delta-logics by a finite set of first-order variables $\Delta = \{v_1, \dots, v_n\}$. Delta-logic formulas are Boolean combinations of *contextual* formulae and *delta-specific* formulae. We now define the former.

Contextual Formulae. Intuitively, a contextual formula $\varphi(\vec{x})$ is a formula that evaluates on a model M while *ignoring* the functions/relations on the locations interpreted for the variables in Δ .

More precisely, a semantic definition of the contextual logic over Σ with respect to Δ is as below. First, we shall define when a pair of models over the same universe and interpretation 'differ only on Δ '.

Definition 3.1 (Models differing only on Δ). Let M and M' be two Σ -models with universe U that interpret constants the same way, and let I be an interpretation of variables over U . Then we say (M, I) and (M', I) differ only on Δ if:

- for every function symbol f and for every $l \in U$, $\llbracket f \rrbracket_M(l) \neq \llbracket f \rrbracket_{M'}(l)$ only if there exists some $v \in \Delta$ such that $I(v) = l$.
- for every relation symbol S and for every $l \in U$, $\llbracket S \rrbracket_M(l) \neq \llbracket S \rrbracket_{M'}(l)$ only if there exists some $v \in \Delta$ such that $I(v) = l$. □

Intuitively, the above says that the interpretation of the (unary) functions and relations of the two models are precisely the same for all elements in the universe that are not interpretations of the variables in Δ .

An $FOL+lf\!p$ formula over Σ , $\varphi(\vec{x})$, is a contextual formula if the formula does not distinguish between models that differ only on Δ :

Definition 3.2 (Contextual formulae). An $FOL+lf\!p$ formula over Σ , $\varphi(\vec{x})$, is a contextual formula if for every two Σ -models M and M' with the same universe and every interpretation I such that (M, I) and (M', I) differ only on Δ , $M, I \models \varphi$ iff $M', I \models \varphi$. □

Contextual formulae can be easily written using syntactic restrictions where the formulae (and the recursive definitions) are written so that every occurrence of $f(t)$ (where f is a function symbol) or $P(t)$ (where P is a relation symbol), where t is a term of type σ_0 , is guarded by the clause " $t \notin \Delta$ ", which is short for $\bigwedge_{v \in \Delta} t \neq v$.

Let us now give some examples of contextual formulae.

*Example: Let us fix a finite set Δ of first-order variables.
The recursive definition*

$$ls^*(x) :=_{lfp} (x = nil \vee (x \neq nil \wedge x \notin \Delta \wedge ls^*(n(x))) \vee (x \neq nil \wedge \bigvee_{v \in \Delta} (x = v \wedge b_v)))$$

is a contextual formula/definition with respect to Δ . The above defines lists but by “imbibing” facts about whether a location v in Δ is a list using the free Boolean variable b_v . These Boolean variables play the role of the interface variables/parameters mentioned earlier. The least fixpoint semantics of the above definition gives a unique definition: $ls(u)$ is true iff there is a path using the pointer $n()$ that either ends in nil or ends in a node v in Δ where b_v is true. Note that the formula $n(x)$ is guarded by the check $x \in \Delta$, and hence is a contextual formula. Changing the model to reinterpret $n()$ over Δ (but preserving the interpretation of the variables $b_v, v \in \Delta$) will not change the definition of ls .

Delta-specific formulae: A delta-specific formula is a quantifier-free formula where every occurrence of a term of the form $f(t)$, (for an uninterpreted function f) $t \in \Delta$. Furthermore the formula does not refer to any of the recursively defined functions/predicates.

Delta-logic formulae: A delta-logic formula is a Boolean combination of contextual formulae and delta specific-formulae.

Example: The formula $u' = n(u) \wedge ls^(u')$ is a delta-logic formula that uses the above definition of ls .*

4 A SEPARABILITY THEOREM

In this section, we show a key result: that for any quantifier-free $FO+lfp$ formula we can effectively find an equivalent quantifier-free delta-logic formula. We do this by separately reasoning with the elements of the formula that are specific to Δ , and those that are oblivious to Δ . We bring these separate analyses together with a set of parameters ‘ P ’ that we shall describe in parts and justify below.

First, let us consider a recursively defined function R . Let the set of functions $PF = \{p_i \mid 1 \leq i \leq k\}$ (for some k) model the pointer fields (we assume that have a clause $p_i(nil) = nil$ for every $1 \leq i \leq k$). We also assume that Δ is fixed for this discussion.

We define a set of variables $\{R^d \mid d \in \Delta\}$ of the type of the range of R . These variables are part of the set of parameters P . Then if R is defined as $R(x) :=_{lfp} \varphi(x)$ we define a new function corresponding to R , namely R^P , that is recursively defined as follows:

$$\begin{aligned} R^P(x) &:=_{lfp} & R^d \text{ if } \llbracket x \rrbracket = \llbracket d \rrbracket \text{ for some } d \in \Delta & \quad (\text{delta case}) \\ & & \varphi[R^P/R] \text{ if } \llbracket x \rrbracket \notin \llbracket \Delta \rrbracket & \quad (\text{recursive case}) \end{aligned}$$

It is easy to see that for a formula $R(x)$, $R^P(x)$ would be a contextual formula since any model of it would not depend on the valuation of PF over Δ (see definition of ls^* in Section 3 for example).

To capture the semantics of the original lfp definition, we constrain these parameters. This will yield a definition that is equivalent in $FO+lfp$ under such constraints.

We do this by writing constraints that, effectively, unfold the recursive definition over Δ . However in doing so we would run into a problem with cycles; for instance, simply imbibing the value of $lseg_z$ from the node pointed to would not work when we have a circular list.

We handle this by introducing the notion of the ‘rank’ of a location w.r.t R . In particular, for the example of a circular list, if we recursively defined rank as a natural number increasing on a list starting from 0 at the location nil , there is no way to provide a valuation of every element on the cycle as pointing to a list. However, since the rank will need to communicate through the elements outside Δ to maintain this order (pointer paths between elements interpreted in Δ need not lie within it), it will also be a similarly relativised lfp definition with its own parameters

$\{RANK_R^d \mid d \in \Delta\}$ which are also included in the parameter set P . We choose to model the rank as a function to $\mathbb{N} \cup \{\perp\}$ (\perp signifies undefined rank) as follows for the recursively defined function R :

$$\begin{aligned} Rank_R(x) &:= lfp \ RANK_R^d & \text{if } \llbracket x \rrbracket = \llbracket d \rrbracket \text{ for some } d \in \Delta (\text{delta case}) \\ 0 & & \text{if } \llbracket x \rrbracket \notin \llbracket \Delta \rrbracket \wedge \varphi(x)[\perp/R] \neq \perp (\text{base case}) \\ \max_{1 \leq i \leq k} \{Rank_R(p_i(x))\} & & \text{if } \llbracket x \rrbracket \notin \llbracket \Delta \rrbracket \wedge R^P(x) \neq \perp (\text{recursive case}) \end{aligned}$$

Intuitively, $Rank_R(x)$ imbibes the value $RANK_R^d$ for x in Δ , and otherwise propagates the rank across the context. This will be used below to infer correctly the value of R on elements in Δ .

Finally, we define a delta-logic formula β_R that ‘computes’ the value of R on Δ by unfolding the definition of R using the pointer fields on Δ . To this end, we define the final sets of parameters to be included in P in similar vein as above to ‘communicate’ from the context to Δ . These parameters will correspond to *boundary* variables: $p_i(\Delta) \setminus \Delta$ for some i , and are therefore named thus:

Let $R^{p_i(\Delta)} = \{R^{p_i(d)} \mid d \in \Delta\}$ of the type of range of R and $RANK_R^{p_i(\Delta)} = \{RANK_R^{p_i(d)} \mid d \in \Delta\} \subseteq \mathbb{N} \cup \{\perp\}$ for every $1 \leq i \leq k$.

We then denote the substitution $\varphi(x)[P/R]$ as replacing the term $R(p_i(x))$ with $R^{p_i(x)}$ for every $1 \leq i \leq k$, and $\varphi(x)[\perp/R]$ as replacing with \perp . With the above, we write the following delta-specific constraint β_R for a recursively defined function R :

$$\begin{aligned} \bigwedge_{d \in \Delta} & \left[\left(\varphi(d)[\perp/R] \neq \perp \implies R^d = \varphi(d)[\perp/R] \wedge (RANK_R^d = 0) \right) \right] (\text{base case}) \\ & \wedge \left(\varphi(d)[\perp/R] = \perp \wedge \varphi(d)[P/R] \neq \perp \implies \left(R^d = \varphi(d)[P/R] \right. \right. \\ & \quad \left. \left. \wedge \left(RANK_R^d = \max_{1 \leq i \leq k} (\{RANK_R^{p_i(d)}\}) + 1 \right) \right) \right) (\text{recursive case}) \\ & \wedge \left(\varphi(d)[\perp/R] = \perp \wedge \varphi(d)[P/R] = \perp \implies \left(R^d = \perp \wedge (RANK_R^d = \perp) \right) \right)] (\text{undefined}) \end{aligned}$$

The above constraints capture accurately the values of R on Δ :

- the base case simply constrains the parameter R^d at a node interpreting its corresponding variable to be the value provided by the recursive function definition R , and its rank to be 0 when the interpretation for that variables satisfies the base case of the recursive definition.
- the recursive case constrains the parameter (when it does not satisfy the base case) to be the value computed by one unfolding of the definition, where the values of the descendants are also denoted by their respective parameters (whether Δ or boundary) and its rank to be one more than the maximum rank among its descendants.
- the undefined case constrains the parameter to be undefined when it must be according to an unfolding of the definition, and its rank to be undefined as well.

Lastly, we must also have that the boundary variables do in fact communicate the values of the contextual recursive definition to Δ , i.e that the placeholder parameters for their values are indeed the values provided by the contextual lfp definition R^P :

$\bigwedge_{1 \leq i \leq k} [p_i(d) \notin \Delta \implies (R^{p_i(d)} = R^P(p_i(d))) \wedge (RANK_R^{p_i(d)} = Rank_R(p_i(d)))]$ This formula in conjunction with the above constraints forms the delta-logic formula β_R introduced above.

Before stating the main theorem of this section, we prove a technical lemma. This lemma states that for any recursively defined function R and corresponding set of parameters P_R (a) there is always a valuation of the parameters P_R that satisfies the constraints above, and (b) any valuation of

the parameters that satisfies the constraints above will make the contextual definition R^{P_R} precisely the same as R .

LEMMA 4.1. *For any recursively defined function R , $(\exists P_R. \beta_R) \wedge (\forall P_R. (\beta_R \implies R^{P_R} = R))$*

We now can show that any quantifier-free FO+*lfp* formula, say α , has an equivalent delta-logic formula. Let the set of recursive functions/predicates mentioned in α be \mathcal{R} , and Δ be fixed. Let $\mathcal{R}^P = \{R^{P_R} \mid R \in \mathcal{R}\}$, and $P_{\mathcal{R}} = \bigcup_{R \in \mathcal{R}} P_R$. Then:

THEOREM 4.2 (SEPARABILITY). $\alpha \equiv \exists P_{\mathcal{R}}. \alpha[\mathcal{R}^P/\mathcal{R}] \wedge \left(\bigwedge_{R \in \mathcal{R}} \beta_R \right)$.

PROOF. Consider that α holds. From Lemma 4.1, for every $R \in \mathcal{R}$, we can pick a valuation for P_R such that β_R holds and therefore, $R^{P_R} = R$. Thus we have that $\alpha[\mathcal{R}^P/\mathcal{R}] \wedge \left(\bigwedge_{R \in \mathcal{R}} \beta_R \right)$ holds.

Conversely, let $\alpha[\mathcal{R}^P/\mathcal{R}] \wedge \left(\bigwedge_{R \in \mathcal{R}} \beta_R \right)$ hold. Again, from Lemma 4.1 we have that for every $R \in \mathcal{R}$, the valuation given by the model for P_R satisfies β_R , and therefore $R^{P_R} = R$. Therefore, $\alpha[\mathcal{R}^P/\mathcal{R}][\mathcal{R}/\mathcal{R}^P] = \alpha$ holds. \square

Observe that the latter formula in Theorem 4.2 is a formula in delta-logic, i.e, is a Boolean combination of contextual formulae and delta-specific formulae.

5 TRANSLATING VERIFICATION CONDITIONS TO DELTA LOGICS

We now explain the crux of the motivation behind delta-logics, namely that they can naturally express verification conditions, without introducing extra quantification or complex reformulation of recursive definitions.

Consider a Hoare Triple: $\langle \alpha_{pre}(X, P, \mathcal{R}), S, \alpha_{post}(X, P, \mathcal{R}) \rangle$, such that α_{pre} and α_{post} utilise a set of relations and functions \mathcal{R} with recursive definitions defined using FO+*lfp*. The program manipulates a set of scalar variables X , and pointer and data fields P , the latter being modeled as unary functions.

The verification condition is then of the form $\alpha_{pre}(X, P, \mathcal{R}) \wedge T(X, X', P, P') \implies \alpha_{post}(X', P', \mathcal{R}')$, where T captures the semantics of the program snippet S , which has no function calls. T can be expressed as a delta-specific formula T_{Δ} conjoined with the formula $\bigwedge_{f \in P} (\forall z. z \notin \Delta \implies f'(z) = f(z))$, which captures the fact that the fields of elements outside Δ have not changed. Notice that the definitions of \mathcal{R}' are obtained from the definitions of \mathcal{R} by substituting P' for P .

Let us now show how to express this VC using an equivalent delta-logic formula without introducing universal quantification.

First, using the separability theorem, Theorem 4.2, we can write α_{pre} and α_{post} as delta-logic formulae $\alpha_{pre}^*(U, X, P, \mathcal{R}^U)$ and $\alpha_{post}^*(V, X', P', (\mathcal{R}')^V)$. Notice that the definition of R' uses the transformed fields P' and we translate using *different* sets of parameters: U and V . However, observe that since $(R')^V$ is a contextual definition, we know that it does not refer to the changed heaplet Δ . Consequently, we can replace P' with P uniformly in $(R')^V$, which yields $(R)^V$. Also, we can *remove* the universally quantified conjunct in T that says that fields for locations outside Δ are unaltered.

This leaves us with a verification condition of the form:

$$\alpha_{pre}^*(U, X, P, \mathcal{R}^U) \wedge T_{\Delta}(X, X', P, P') \implies \alpha_{post}^*(V, X', P, (\mathcal{R})^V)$$

where T_{Δ} is a delta-specific formula that simply says how the variables X change to X' , and the fields P change to P' over Δ . The VC is now clearly a delta-logic formula.

Notice that though we have two set of recursive definitions, the definitions do not rely on different sets of data and pointer fields (which are functions), but rather differ only on the first-order parameter variables U and V .

6 A DECIDABLE DELTA LOGIC ON LISTS WITH LIST MEASURES

In this section, we will define a delta-logic on linked lists equipped with *list measures*— measures of list segments that include its length, heaplet, the multiset of keys stored in the list (say, in a data-field key), and the minimum and maximum keys stored in it. We prove that the quantifier-free first-order logic fragment of this delta-logic is *decidable*.

We prove decidability by first proving that the corresponding contextual logic formulae can be translated to equisatisfiable quantifier-free first-order formulae. Consequently, a delta-logic formula, being a Boolean combination of contextual formulae and delta-specific formulae can be translated to quantifier-free formulae as well, which is decidable using a Nelson-Oppen combination of decidable theories of arithmetic, sets, and uninterpreted functions.

The contextual logic of list measures

We now define a contextual logic over list segments and measures over them. Notice from Section 4 that the separation of formulae into delta-logic introduces recursive definitions in the contextual logic (parameterized over various sets of variables) and additionally a rank function for each such definition. However, it is easy to see that for the definitions of lists and measures, all the rank functions (over a given set of parameters) coincide, since the existence of a meaningful value for each of these measures is predicated upon the referred location pointing to a list. This motivates the following recursive definitions for our contextual logic of lists and measures.

As usual, let us fix a set of first-order variables Δ . Let us also fix a single pointer field n .

Definition 6.1 (Recursive Definitions for the Logic of List Measures (LM)). Let us fix a set of *parameter variables* P that consists of the following of sets of variables: a set of Boolean variables LS_z^v , a set of variables with type set of locations HLS_z^v , a set of variables of type multiset of keys $MSKEYS_z^v$, and a set of variables of type integer (type of keys in general) Max_z^v and Min_z^v , where z, v range over Δ .

The contextual logic of list measures (LM) wrt Δ and parameter variables P is defined using the following recursive definitions, which depend crucially on the parameter variables P .

- We have unary relations ls_z^P that capture linked list segments that end in z , where the relation for a location v in Δ is imbibed using the Boolean variables LS_z^v ($v \in \Delta$), and where z is any element of Δ or the constant location nil . (The relation $ls_{nil}()$ captures whether a location points to a list ending with nil .)

This is defined as follows:

$$ls_z^P(x) :=_{lfp} \left(x = z \vee \left(x \neq z \wedge x \neq nil \wedge x \notin \Delta \wedge ls_z^P(n(x)) \right) \vee \left(x \neq z \wedge x \in \Delta \wedge \bigwedge_{v \in \Delta} (x = v \Rightarrow LS_z^v) \right) \right)$$

- We have recursive definitions that capture the heaplet of such list-segments, where the heaplet of list-segments from an element v in Δ to z (where $z \in \Delta \cup \{nil\}$) is imbibed from the set variable HLS_z^v :

$$\begin{aligned} hls_z^P(x) &:=_{lfp} && \emptyset \text{ if } \llbracket x \rrbracket = \llbracket z \rrbracket \\ &&& \{x\} \cup hls_z^P(n(x)) \text{ if } \llbracket x \rrbracket \neq \llbracket z \rrbracket \wedge \llbracket x \rrbracket \neq \llbracket nil \rrbracket \wedge \llbracket x \rrbracket \notin \llbracket \Delta \rrbracket \\ &&& HLS_z^v \text{ if } \llbracket x \rrbracket \neq \llbracket z \rrbracket \wedge \llbracket x \rrbracket = \llbracket v \rrbracket \wedge v \in \Delta \end{aligned}$$

$$\begin{aligned}
\text{Location Term } lt &::= x \mid p_i(y) \mid \text{nil where } y \notin \Delta \\
\text{Integer Term } it &::= c \mid \text{len}_z^P(lt) \mid it + it \\
\text{Key Term } keyt &::= c \mid \text{key}(lt) \mid \text{max}_z^P(lt) \mid \text{min}_z^P(lt) \mid keyt + keyt \\
\text{Heaplet Term } hlt &::= \emptyset \mid \{lt\} \mid \text{hls}_z^P(lt) \mid hlt \cup hlt \mid hlt \cap hlt \mid hlt \setminus hlt \\
\text{MultisetKeys Term } mskt &::= \emptyset \mid \text{mskeys}_z^P(x) \mid mskt \cup_m mskt \mid mskt \cap_m mskt \mid mskt \setminus_m mskt \\
\\
\text{Formulas } \varphi &::= \text{true} \mid \text{false} \mid \text{ls}_z^P(x) \mid \text{sorted}_z^P(x) \mid lt = lt \mid lt \in hlt \mid hlt \subseteq hlt \mid hlt = \emptyset \mid \\
&it < it \mid it = it \mid keyt < keyt \mid keyt = keyt \mid \\
&keyt \in mskt \mid mskt \subseteq_m mskt \mid \varphi \vee \varphi \mid \neg \varphi
\end{aligned}$$

Fig. 1. The context logic *LM* of list measures involving list-segments, heaplets, multisets of keys, max, min, and sortedness.

- We have similar recursive definitions that capture the multiset of data elements stored in list segments, the maximum/minimum element stored in the list, and a predicate capturing sortedness of list segments. We omit these definitions: see Appendix for detailed definitions.

We define the *contextual logic of list-measures (LM)* to be quantifier-free formulae that use only the recursive definitions of *LM* mentioned above, and combine them as described in Figure 1. The logic *LM* allows first-order variables that range over locations, keys, and integers. Note that dereferencing pointers of locations is completely disallowed— the delta-specific formula will refer to such dereferences, depending on the particular state modeled. Here Δ is assumed to be a subset of the free location variables in the contextual formula. Formulae in *LM* are allowed to refer to recursive predicates/functions defined using over various sets of parameter variables.

6.1 Translating contextual formulae to quantifier-free recursion-free formulae

We can now state the main result of this section:

THEOREM 6.2. *For any quantifier-free formula $\varphi(\mathcal{P}, X)$ of $LM[ls, hls, rank, len, min, max, sorted]$, the quantifier-free and recursion-free formula $\psi(\mathcal{P}, X)$ obtained from the translation below satisfies the following property: for any interpretation of the free variables in $\mathcal{P} \cup X$, there is a model for φ iff there is a model for ψ .*

COROLLARY 6.3. *The delta-logic of linear measures is decidable.*

Let us fix a set of sets of parameters $\mathcal{P} = \{P_1, \dots, P_k\}$ (we encourage the reader to fix $k = 2$ in their mind while reading the section, as it's the most common and the logic VCs translate to, as shown in Section 5).

We will first describe the decision procedure and its proof of correctness for the fragment of *LM* that involves only the three recursive definitions ls_z^P , hls_z^P , and rank_z^P , where $P \in \{\mathcal{P}\}$, which we will refer to as $LM[ls, hls, rank]$. Then we will extend the procedure to handle the logic with all the other measures; this latter proof requires more expressive decision procedures and pseudo-measures that make its proof harder.

Let us assume a quantifier-free $LM[ls, hls, rank]$ formula φ which is a Δ -logic formula w.r.t a finite set of variables Δ . Assume the set of (free) location variables occurring in φ is $X = \{x_1, \dots, x_n\}$ with $\Delta \subseteq X$.

In order to determine whether there is a model satisfying φ , we need to construct a universe of locations, an interpretation of the variables in X , and the heap (with the single pointer field $n()$) on all locations *outside* Δ (the definition of $n()$ on Δ , by definition, does not matter).

Our decision procedure relies intuitively on the following observations. First, note that the locations reached by using the $n()$ pointer any number of times forms the relevant set of locations that φ 's truth can depend on (as φ is quantifier-free and has recursive definitions that only use the n -pointer). There are three distinct cases to consider when pursuing the paths using the n -pointer on a location x : (a) the path may reach a node in Δ , (b) the path may reach a node that is reachable also from another location in X , or (c) the path may never reach a location in Δ nor a location that is reachable from another location in X .

The key idea is to *collapse* paths where the reachability of those locations from locations in (interpreted by) X does not change. More precisely, let L be the set of all locations reachable from X such that l is in Δ or for every location l' reachable from X such that $n(l') = l$, the set of nodes in X that have a path to l' is different from the set of nodes in X that have a path to l .

It is easy to see that there are at most $|X| - 1$ locations of the above kind that are distinct from Δ , since the paths can merge at most $|X| - 1$ times forming a tree-like structure. Our key idea is now to represent these list segments that connect these kinds of locations *symbolically*, summarising the measures on these list segments. Since there are only a bounded number of such locations and hence list segments, we can compute recursive definitions of linear measures involving them using quantifier-free and recursive-definition-free formulae.

We construct a formula ψ that is satisfiable iff φ is satisfiable, as follows. First, we fix a new set (distinct from X) of location variables $V = v_1, \dots, v_{|X|-1}$, to stand for the merging locations described above. We introduce an uninterpreted function $T : V \cup (X \setminus \Delta) \longrightarrow V \cup X \cup \{\perp\}$ (\perp is used to signify that the $n()$ -path on the location never intersects $X \cup L$). Let Z be the set of variables in X such that the recursive definitions $ls_z^P, hls_z^P, rank_z^P$, for some $P \in \mathcal{P}$, occur in φ .

ψ is the conjunct of the following formulae:

- The formula φ (but with recursive definitions treated as uninterpreted relations and functions).
- For every $z \in Z$, we introduce an uninterpreted function $Dist_z : V \cup (X \setminus \Delta) \longrightarrow \mathbb{N} \cup \{\perp\}$ that is meant to capture the distance from any location in $V \cup X$ to z , if z is reachable from that location without going through Δ , and is \perp otherwise. We add the constraint:

$$\begin{aligned} & \bigwedge_{v \in V \cup (X \setminus \Delta)} [(Dist_z(v) = 0 \Leftrightarrow v = z) \wedge \\ & v \neq z \Rightarrow ((T(v) = \perp \vee Dist_z(T(v)) = \perp) \Rightarrow Dist_z(v) = \perp) \\ & \wedge ((T(v) \neq \perp \wedge Dist_z(T(v)) \neq \perp) \Rightarrow Dist_z(v) = Dist_z(T(v)) + 1))] \end{aligned}$$

- For every $x \in X$, and for every $P \in \mathcal{P}$, we have a conjunct:

$$ls_z^P(x) \Leftrightarrow (Dist_z(x) \neq \perp \vee \bigvee_{v \in \Delta} (Dist_v(x) \neq \perp \wedge LS_z^v))$$

- For every $x \in X$, $z \in Z$, and for every $P \in \mathcal{P}$, we have a conjunct:

$$(Dist_z(x) = \perp \Rightarrow rank_z^P(x) = \perp) \wedge (Dist_z(x) \neq \perp \Rightarrow rank_z^P(x) = RANK_z^P)$$

- We capture the heaplets of list-segments from $v \in V \cup (X \setminus \Delta)$ to $T(v)$ (excluding both end-points) using a set of locations $H(v)$ and constrain them so that they are pairwise disjoint and do not contain the locations X :

$$\bigwedge_{x \in X, v \in V \cup (X \setminus \Delta)} x \notin H(v) \wedge \bigwedge_{v, v' \in V \cup (X \setminus \Delta)} (v \neq v' \Rightarrow H(v) \cap H(v') = \emptyset)$$

- We can then precisely capture the heaplet $hls_z^P(x)$ by taking the union of all heaplets of list segments lying on its path to z . We do this using the following constraint, for each $v \in X \cup V$:

$$\begin{aligned} (Dist_z(v) = \perp \Rightarrow hls_z^P(v) = \emptyset) \wedge (hls_z^P(z) = \emptyset) \wedge \\ (Dist_z(v) \neq \perp \wedge v \neq z) \Rightarrow hls_z(v) = H(v) \cup \{v\} \cup hls_z(T(v)) \end{aligned}$$

Note that the formula ψ is quantifier-free and over the combined theory of arithmetic, uninterpreted functions, and sets.

We can show the correctness of the above translation:

THEOREM 6.4. *The statement of Theorem 6.2 holds for the fragment $LM[ls, hls, rank]$.*

We now turn to the more complex logic $LM[ls, hls, rank, len, mskeys, min, max, sorted]$, and show that any quantifier-free formula φ in the logic can be similarly translated. First, we model the multiset of keys, minimum and maximum values and sortedness of each list-segment from v to $T(v)$ (where $v \in (X \setminus \Delta) \cup V$), which is outside Δ , using multiset variables $mskeys\mu(v)$, integer variables $min\mu(v)$, $max\mu(v)$, and $len\mu(v)$ and boolean variables $sorted\mu(v)$. We can also aggregate them, as above, to express the sets $mskeys_z(x)$, $min_z(x)$, $max_z(x)$, $len_z(x)$ and $sorted_z(x)$, for each $z \in Z$ and each $x \in (X \setminus \Delta) \cup V$, similar to definitions of $hls_z(x)$ as defined above. A point of note is that the recursive definition of sortedness across segments is expressed by using both $Min_z(x)$ and $Max_z(x)$ definitions, though the recursive definition of sortedness uses only minimum— this is needed as expressing when the concatenation of sorted list segments is sorted requires the max value of the first segment. We skip these definitions as they are easy to derive.

The main problem that remains is in *constraining* these measures so that they can be the measures of the *same* list segment, i.e, be the attributes not of a pseudo-model. The following constraints capture this, for each $v \in (X \setminus \Delta) \cup V$:

- The cardinality of $hls\mu(v)$ must be $len\mu(v)$.
- The cardinality of $mskeys\mu(v)$ must be $len\mu(v)$.
- $min\mu(v)$ and $max\mu(v)$ must be the minimum and maximum elements of $mskeys\mu(v)$.
- If $min\mu(v) = max\mu(v) \neq \perp$, then $sorted\mu(v)$ can only be true.

The intuition is that any measures meeting the above constraints can be realized using true list segments. As for the fourth clause above, notice that any list segment with minimum element different from maximum can be realized by either a sorted or an unsorted list.

The above constraints on measures, though seemingly simple, are hard to shoehorn into existing decidable theories (though the fourth constraint can be easily expressed). The first two constraints can be expressed using quantifier-free BAPA [7] (Boolean Algebra with Presburger Arithmetic) constraints, which is decidable. We can get around defining the minimum of list segments by having the set of keys store only offsets from the minimum (and including the key 0 always). However, capturing max and sortedness measures as well while preserving decidability seems hard.

Consequently, we give a new decision procedure that exploits the setup we have here. First, note that we can restrict the formulae that use sets containing keys to involve only membership testing of free variables in them, combinations using union and intersection, and checking emptiness of derived sets. We can *disallow checking non-emptiness* as non-emptiness of a set S can always be captured by demanding $k \in S$, for a fresh free variable k .

Our primary observation is that we can then restrict the multiset of keys to be over a *bounded* universe of elements. This bounded universe consists of one element for each free variable of type key in the formula (call this K), and, in addition, will consist of one element for each Venn region formed by the multiset of keys for each segment $(v, T(v))$ of the context's heap. The idea

Correct programs	#VCs	time(s)	Buggy programs	time(s)
append(x: list, y: list)	4	57	buggy_append(x: list, y: list)	0.2
copyall(x: list)	4	183	buggy_copyall(x: list)	36
detect_cycle(x: list)	6	5	buggy_detect_cycle(x: list)	0.7
deleteall(x: list, k: key)	5	10	buggy_deleteall(x: list, k: key)	0.1
find(x: list, k: key)	3	6	buggy_find(x: list, k: key)	0.1
insert(x: list, k: key)	4	38	buggy_insert(x: list, k: key)	0.2
insert_front(x: list, k: key)	1	3	buggy_insert_front(x: list, k: key)	0.1
insert_back(x: list, k: key)	4	20	buggy_insert_back(x: list, k: key)	0.2
reverse(x: list)	3	14	buggy_reverse(x: list)	0.1
sorted_append(x: list, y: list)	4	54	buggy_sorted_append(x: list, y: list)	0.2
sorted_deleteall(x: list, k: key)	5	2	buggy_sorted_deleteall(x: list, k: key)	1.1
sorted_insert(x: list, k: key)	4	17	buggy_sorted_insert(x: list, k: key)	1.2
sorted_reverse(x: list)	3	37	buggy_sorted_reverse(x: list)	0.5
sorted_merge(x: list, y: list)	8	300	buggy_sorted_merge(x: list, y: list)	0.7
insert_back_rec(x: list, k: key)	2	1.1	buggy_insert_back_rec(x: list, k: key)	0.2
deleteall_rec(x: list, k: key)	3	1.3	buggy_deleteall_rec(x: list, k: key)	0.25
even_split_rec(x: list)	2	0.3	buggy_even_split_rec(x: list)	0.45
sorted_merge_rec(x: list, y: list)	4	0.5	buggy_sorted_merge_rec(x: list, y: list)	0.43

Fig. 2. Experimental results for the decision procedure for the delta-logic LM; extended with frame reasoning for function calls

of introducing an element for each Venn region is not new, and is found in many works that deal with combinations of sets and cardinality constraints [7].

Once we have bounded the universe of keys, we can represent a multiset of keys using a set of natural numbers that represent the multiplicity of elements, and write the effect of unions and intersections using Presburger arithmetic. The cardinality of the multiset is the sum of these numbers and the minimum and maximum key can be expressed using the smallest and largest keys in the finite universe with multiplicity greater than 0. We can hence translate the formula into a quantifier-free formula, giving the required theorem. This proves Theorem 6.2.

Note that the above procedure introduces an exponential number of variables, and hence poses challenges to be effective in practice. There are several possible ways of mitigating this. First, there is existing work (see [8]) on reasoning with BAPA that argues and builds practical algorithms that introduce far smaller universes in practice. Second, in the case where we allow only combinations of sets using union (and not intersection), and allow checking subset constraints and emptiness, we can show that introducing a *single* new element in the universe other than K suffices. The reason is that without intersections, the identity of the elements do not matter and their multiplicities are preserved by representing with only one element. We exploit this in our implementation below.

6.2 Implementation and evaluation of decision procedure for LM

We implemented the decision procedure for the delta-logic LM[*ls, hls, rank, len, mskeys, min, max, sorted*] using the reduction to SMT described above, solved using Z3. We applied our technique to a suite of list-manipulating programs.

The specifications for the programs were strong. For example, programs such as `deleteall` which removed a certain key from a list were also verified with the additional requirement any other arbitrary key had to be preserved in multiplicity across the program. We could also verify such properties as the heaplet of the resulting list being a subset of the original heaplet. The `detect_cycle` program was able to express the existence of cycles using the predicates in *LM*, and used them to verify Floyd’s tortoise and the hare algorithm.

The results are summarized in Figure 2. The left column shows results of verifying correct programs, while the column on the right shows results on ‘buggy’ variants of these programs. The buggy programs were obtained by expressing weaker annotations. Our tool worked well on all these examples and for the buggy programs, the satisfying valuation it provided was sufficiently informative to diagnose the error. To the best of our knowledge, ours is the only decidable verification tool that can handle this suite of examples.

The experiments were performed on a machine with an Intel® Core™i7-7600U processor with clock speeds upto 3.9GHz.

7 RELATED WORK

There is rich literature on decidable fragments of separation logic [?] with inductive predicates. While the first works such as [1, 2, 5, 12] handled only list segments, there has been further improvement. The work in [13] handles only list segments, but allows conjunctions with pure formulae from arbitrary decidable SMT theories. The work in [3] argues decidability for separation logic with arbitrary/user-defined inductive predicates, and [9, 10] adapts the idea of using equisatisfiable ‘base’ formulas instead of inductive predicates to include arithmetic constraints satisfying a certain property. However, none of these would be directly applicable to program verification since they do not address entailment, and do not contain an implication in their logic. The work in [13] does contain implications, but it does not allow for spatial formulas other than points-to, list segment and in particular not provide support for multiple natural measures which our work provides. However, in the light of our broader work on context-logics and delta-logics, it is valuable to consider the satisfiability problem in the context-logic since it models the unchanging heap and does not need to address entailments, of which these works above provide useful decision procedures.

In contrast, works such as [14, 15] do address program verification and provide a decidable logic on spatial formulas that when reduced to SMT can be extended with other decidable SMT theories. There is reasonable overlap in the spirit of this claim to our work, where allowing expressions in arbitrary decidable SMT theories on Δ would preserve decidability as well, and would fit well within the paradigm of delta-logics. But where [14] can support trees/tree segments, it does not provide support for measures. Works such as [6, 11] differ similarly. Lastly there is work in the spirit of [4, 16], where one dispenses with decidability and is therefore incomparable in that respect with our work. In particular, [16] is a powerful method capable of verifying all of our examples, but such logics are incomplete and cannot give meaningful counterexamples like our work does.

REFERENCES

- [1] Josh "Berdine, Cristiano Calcagno, and Peter W." O'Hearn. "2004". "A Decidable Fragment of Separation Logic". In *"Proceedings of the 24th International Conference on Foundations of Software Technology and Theoretical Computer Science" ("FSTTCS'04")*. "97–109".
- [2] Josh "Berdine, Cristiano Calcagno, and Peter W." O'Hearn. "2005". "Symbolic Execution with Separation Logic". In *"Proceedings of the Third Asian Conference on Programming Languages and Systems" ("APLAS'05")*. "52–68".
- [3] James "Brotherston, Carsten Fuhs, Juan A. Navarro Pérez, and Nikos" Gorogiannis. "2014". "A Decision Procedure for Satisfiability in Separation Logic with Inductive Predicates". In *"Proceedings of the Joint Meeting of the Twenty-Third EACSL Annual Conference on Computer Science Logic (CSL) and the Twenty-Ninth Annual ACM/IEEE Symposium on Logic in Computer Science (LICS)" ("CSL-LICS '14")*. "25:1–25:10".
- [4] "W. N. Chin, C. David, H. H. Nguyen, and S. Qin". "2007". "Automated Verification of Shape, Size and Bag Properties". In *"12th IEEE International Conference on Engineering Complex Computer Systems (ICECCS 2007)"*. "307–320".
- [5] Byron "Cook, Christoph Haase, Joël Ouaknine, Matthew Parkinson, and James" Worrell. "2011". "Tractable Reasoning in a Fragment of Separation Logic". In *"Proceedings of the 22Nd International Conference on Concurrency Theory" ("CONCUR'11")*. "235–249".
- [6] Radu "Iosif, Adam Rogalewicz, and Jiri" Simacek. "2013". "The Tree Width of Separation Logic with Recursive Definitions". In *"Proceedings of the 24th International Conference on Automated Deduction" ("CADE'13")*. "21–38".
- [7] Viktor "Kuncak, Huu Hai Nguyen, and Martin" Rinard. "2005". "An Algorithm for Deciding BAPA: Boolean Algebra with Presburger Arithmetic". In *"Proceedings of the 20th International Conference on Automated Deduction" ("CADE' 20")*. "260–277".
- [8] Viktor "Kuncak and Martin" Rinard. "2007". "Towards Efficient Satisfiability Checking for Boolean Algebra with Presburger Arithmetic". In *"Proceedings of the 21st International Conference on Automated Deduction: Automated Deduction" ("CADE-21")*. "215–230".
- [9] Quang Loc "Le, Jun Sun, and Wei-Ngan" Chin. "2016". "Satisfiability Modulo Heap-Based Programs". In *"Computer Aided Verification"*. "382–404".
- [10] Quang Loc Le, Makoto Tatsuta, Jun Sun, and Wei-Ngan Chin. 2017. A Decidable Fragment in Separation Logic with Inductive Predicates and Arithmetic. In *Computer Aided Verification*.
- [11] P. "Madhusudan, Gennaro Parlato, and Xiaokang" Qiu. "2011". "Decidable Logics Combining Heap Structures and Data". In *"Proceedings of the 38th Annual ACM SIGPLAN-SIGACT Symposium on Principles of Programming Languages" ("POPL '11")*. "611–622".
- [12] Juan Antonio "Navarro Pérez and Andrey" Rybalchenko. "2011". "Separation Logic + Superposition Calculus = Heap Theorem Prover". In *"Proceedings of the 32Nd ACM SIGPLAN Conference on Programming Language Design and Implementation" ("PLDI '11")*. "556–566".
- [13] Juan Antonio Navarro Pérez and Andrey Rybalchenko. 2013. Separation Logic Modulo Theories. In *Programming Languages and Systems (APLAS)*. Springer International Publishing, Cham, 90–106.
- [14] Ruzica "Piskac, Thomas Wies, and Damien" Zufferey. "2014". "Automating Separation Logic with Trees and Data". In *"Proceedings of the 16th International Conference on Computer Aided Verification - Volume 8559"*. "711–728".
- [15] Ruzica Piskac, Thomas Wies, and Damien Zufferey. 2014. GRASShopper. In *Tools and Algorithms for the Construction and Analysis of Systems*. 124–139.
- [16] Xiaokang "Qiu, Pranav Garg, Andrei Ștefănescu, and Parthasarathy" Madhusudan. "2013". "Natural Proofs for Structure, Data, and Separation". In *"Proceedings of the 34th ACM SIGPLAN Conference on Programming Language Design and Implementation" ("PLDI '13")*. "231–242".

APPENDIX

We give here the detailed definitions of the other measures that were omitted in Section 4.

- We have recursive definitions that capture the multiset of data elements (through a data-field *key*) stored in list segments, where again the multiset of data of list-segments from an element v in Δ to z (where $z \in \Delta \cup \{\text{nil}\}$) is imbibed from the set variable $MSKeys_z^v$:

$$\begin{aligned} mskeys_z^P(x) &:=_{lfp} && \emptyset \text{ if } \llbracket x \rrbracket = \llbracket z \rrbracket \\ &&& \{key(x)\} \cup_m mskeys_z^P(n(x)) \text{ if } \llbracket x \rrbracket \neq \llbracket z \rrbracket \wedge \llbracket x \rrbracket \neq \llbracket \text{nil} \rrbracket \wedge \llbracket x \rrbracket \notin \llbracket \Delta \rrbracket \\ &&& MSKeys_z^x \text{ if } \llbracket x \rrbracket \neq \llbracket z \rrbracket \wedge \llbracket x \rrbracket = \llbracket v \rrbracket \wedge v \in \Delta \end{aligned}$$

- We have recursive definitions that capture the maximum/minimum element of data elements stored in list segments, where again the maximum/minimum element of list-segments from an element v in Δ to z where $z \in \Delta \cup \{\text{nil}\}$ is imbibed from the data variable Max_z^v (or Min_z^v). We assume the data-domain has a linear-order \leq , and that there are special constants $-\infty$ and $+\infty$ that are the minimum and maximum elements of this order. Let $max(r_1, r_2) \equiv ite(r_1 \leq r_2, r_2, r_1)$.

$$\begin{aligned} Max_z^P(x) &:=_{lfp} && -\infty \text{ if } \llbracket x \rrbracket = \llbracket z \rrbracket \\ &&& max(key(x), Max_z^P(n(x))) \text{ if } \llbracket x \rrbracket \neq \llbracket z \rrbracket \wedge \llbracket x \rrbracket \neq \llbracket \text{nil} \rrbracket \wedge \llbracket x \rrbracket \notin \llbracket \Delta \rrbracket \\ &&& MSKeys_z^v \text{ if } \llbracket x \rrbracket \neq \llbracket z \rrbracket \wedge \llbracket x \rrbracket = \llbracket v \rrbracket \wedge v \in \Delta \end{aligned}$$

The function Min_z^P is similarly defined.

- We have a recursive definition that captures sortedness, using the minimum measure.

$$\begin{aligned} Sorted_z^P(x) &:=_{lfp} \Big(x = z \vee \\ &\quad \left(x \neq z \wedge x \neq \text{nil} \wedge x \notin \Delta \wedge min_z^P(x) \neq \perp \wedge key(x) \leq min_z^P(x) \wedge Sorted_z^P(n(x)) \right) \vee \\ &\quad \left(x \neq z \wedge x \in \Delta \wedge min_z^P(x) \neq \perp \wedge key(x) \leq min_z^P(x) \wedge \bigwedge_{v \in \Delta} (x = v \Rightarrow SORTED_z^v) \right) \Big) \end{aligned}$$